

Energy-Efficient 5G Control via Constrained Reinforcement Learning: A Lagrangian PPO Approach

Mini Report

November 27, 2025

Abstract

This report addresses the challenge of minimizing energy consumption in 5G macro base stations while strictly adhering to Quality of Service (QoS) constraints. We formulate the problem as a Constrained Markov Decision Process (CMDP) and propose a solution using Proximal Policy Optimization (PPO) augmented with the Lagrangian relaxation method. Experimental results demonstrate that our approach significantly outperforms static baselines, achieving a reduction in energy consumption while maintaining zero violations of critical KPIs such as drop rate and latency.

1 Problem Formulation

We target 5G macro scenarios where base stations must optimize the trade-off between power consumption and network performance. Unlike naive "always-on" strategies which are wasteful, or heuristic sleep modes that risk QoS degradation, we aim to learn a dynamic control policy.

The objective is to minimize the cumulative energy e_t while satisfying a set of K Key Performance Indicator (KPI) constraints (e.g., Drop Rate thresholds, Latency limits, CPU/PRB utilization ceilings). This is formally cast as a Constrained Markov Decision Process (CMDP):

$$\min_{\pi_\theta} J_E(\pi) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^T \gamma^t e_t \right] \quad (1)$$

$$\text{s.t.} \quad J_{C_k}(\pi) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^T \gamma^t c_t^{(k)} \right] \leq \delta_k, \quad \forall k \in \{\text{Drop, Latency, CPU, PRB}\} \quad (2)$$

where $c_t^{(k)}$ represents the penalty signal for constraint k derived from the simulator state.

2 Methodology

To solve the CMDP, we employ the **Lagrangian Proximal Policy Optimization (PPO)** framework.

2.1 Why Lagrangian PPO over Simple Reward Shaping?

A common baseline approach is *Simple Reward Shaping* (or static penalty), where constraints are incorporated directly into the reward function: $r_t = -e_t - \beta \sum c_t^{(k)}$. However, this method suffers from critical limitations:

1. **Sensitivity to Hyperparameters:** Determining the optimal fixed weight β is notoriously difficult. If β is too small, the agent ignores constraints to maximize energy savings. If β is too large, the agent becomes overly conservative, potentially failing to learn any meaningful control policy (vanishing gradient for the main objective).
2. **Instability & Sub-optimality:** A static weight alters the optimization landscape permanently. It does not adapt to the changing dynamics of the environment, often leading to policies that oscillate or get stuck in local optima without satisfying constraints.

Our Solution: The Lagrangian relaxation method transforms the constrained problem into an unconstrained min-max game (Primal-Dual optimization). Instead of a fixed β , we introduce learnable Lagrange multipliers λ_k that dynamically adapt based on constraint violations:

$$\max_{\lambda \geq 0} \min_{\theta} L(\theta, \lambda) = J_E(\pi_{\theta}) + \sum_k \lambda_k (J_{C_k}(\pi_{\theta}) - \delta_k) \quad (3)$$

The reward passed to the PPO agent at each timestep becomes $r_t = -e_t - \sum_k \lambda_k c_t^{(k)}$. The multipliers λ are updated via gradient ascent:

$$\lambda_k^{new} = \Pi_{[0, \lambda_{max}]} (\lambda_k + \alpha (J_{C_k} - \delta_k)) \quad (4)$$

This mechanism guarantees that the penalty "tightens" only when constraints are violated and relaxes when safety is assured, allowing the agent to safely explore energy-saving actions.

2.2 Implementation Details

The architectural choices implemented in `app/energy_agent/rl_agent.py` include:

- **State Representation:** A composite vector of global features (7 dims), augmented features (10 dims), and per-cell local features ($8 \times N_{cells}$), normalized and zero-padded to a fixed size.
- **Algorithm:** Gaussian Actor-Critic PPO. The loss function combines the clipped surrogate objective with the Lagrangian cost:

$$L = L^{CLIP}(\theta) - c_1 L^{VF} + c_2 S[\pi] - \sum \lambda_k L_k^{Cost}$$

- **Training:** Distributed experience collection across $n_{env} = 4$ parallel simulators using Generalized Advantage Estimation (GAE).

3 Experimental Setup

We benchmark the proposed controller against standard baselines across public scenarios defined in `app/scenarios`:

1. **Random Policy:** Selects uniform power ratios randomly. This serves as a stress test for QoS stability.
2. **Max Power (Baseline):** Deterministic "always-on" (action 1.0). This represents the upper bound for energy consumption.
3. **Lagrangian PPO (Final RL):** The trained policy checkpoint (`models/final.pth`).

Experiments were conducted using the evaluation script `app/benchmark_models.py`, recording metrics into `reports/benchmark.csv` for analysis.

4 Results and Discussion

The performance trade-offs are visualized in the facet chart (`reports/benchmark_facets.png`).




Figure 1: Comparison of Energy, Drop Rate, and Violations across policies.

Key observations from the quantitative analysis:

- **Energy Efficiency:** The *Final RL* agent consistently reduces energy consumption by **2.4% – 3.0%** compared to the Max Power baseline. While the Random policy occasionally matches these energy levels, it fails to maintain network stability.
- **Constraint Satisfaction (Drop Rate):** The Random policy is unreliable, averaging a drop rate of 0.53–0.82% in heavy traffic scenarios, causing frequent threshold violations (7–18 instances). In contrast, *Final RL* maintains the drop rate tightly between **0.27–0.32%**, incurring **zero violations** across all test cases.

- **Control Smoothness:** Beyond simple metrics, the RL agent demonstrates smoother control, reducing handover counts by $\sim 50\text{--}90\%$ relative to random actions.

5 Conclusion

The proposed Lagrangian PPO controller successfully identifies a practical Pareto optimal point for 5G base stations. By utilizing a primal-dual optimization approach, we overcame the instability inherent in simple reward shaping. The result is a robust policy that delivers consistent energy savings while satisfying strict QoS constraints, making it a viable candidate for real-world deployment compared to heuristic or uncontrolled baselines.